


# Advanced SQL

# Advanced SQL

- ❑ Data Types
- ❑ Integrity Constraints
- ❑ Trigger
- ❑ Authorization
- ❑ Index



# Data Types

---

# Built-in Data Types in SQL

- ❑ **date:** Dates, containing a (4 digit) year, month and date
  - Example: **date** '2005-7-27'
- ❑ **time:** Time of day, in hours, minutes and seconds.
  - Example: **time** '09:00:30'      **time** '09:00:30.75'
- ❑ **timestamp:** date plus time of day
  - Example: **timestamp** '2005-7-27 09:00:30.75'
- ❑ **interval:** period of time
  - Example: **interval** '1' day
  - Adding or subtracting an interval to a date or time gives back a date or time, respectively

# Build-in Data Types in SQL (Cont.)

- ❑ Can extract values of individual fields from date/time/timestamp
  - **extract** (*field* from *d*)
    - ▶ *field*: can be one of year, month, day, hour, minute, or second.
    - ▶ *d*: date or time value
  - Example: `SELECT EXTRACT(YEAR FROM '2019-09-22')`
- ❑ Can cast string types to date/time/timestamp
  - Example: **cast** <string-valued-expression> **as date**
  - Example: **cast** <string-valued-expression> **as time**

# User-Defined Types

- ❑ **create type** creates user-defined type in SQL

**create type** *Dollars* **as numeric (12,2) final**

- **final**: no meaning here, required by the SQL:1999

**create table** department

(dept\_name **varchar** (20),

building **varchar** (15),

budget *Dollars*);

- ❑ **create domain** construct in SQL-92 creates similar user-defined domain types

**create domain** *person\_name* **char(20) not null**

# User-Defined Types

- ❑ User-defined types and domains are similar.
- ❑ Differences between types and domains
  - Domains can have constraints, such as **not null** specified on them, and can have default values defined for variables of the domain type
  - Domains are **not strongly typed**. Values of one domain type can be assigned to values of another domain type as long as the underlying types are compatible
  - Example: **create domain Dollars numeric(12, 2)**  
**create domain Pounds numeric(12, 2)**


# Domain Constraints

- ❑ **Domain constraints** are the most elementary form of integrity constraint.
  - They test values inserted in the database, and test queries to ensure that the comparisons make sense.
- ❑ We **cannot** assign or compare a value of **type** Dollars to a value of **type** Pounds.
  - However, we can convert type as below  
(**cast** *r.A as Pounds*)  
(Should also multiply by the dollar-to-pound conversion-rate)



# Large-Object Types

- ❑ Large objects (photos, videos, CAD files, etc.) are stored as a *large object*.
    - **blob**: binary large object
      - ▶ object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
    - **clob**: character large object
      - ▶ object is a large collection of character data
- book\_review* **clob**(10KB)
- image* **blob**(10MB)
- movie* **blob**(2GB)
- When a query returns a large object, a pointer is returned rather than the large object itself.



# Integrity Constraints

---

# Integrity Constraints

- ❑ **Integrity Constraints (ICs):** conditions that must be true for **any** instance of the database; e.g., **data type**.
  - An instructor name cannot be null
  - A salary of a bank employee must be at least \$4.00 an hour
  - No two instructors can have the same instructor ID
  - ICs are specified when schema is defined.
  - ICs are checked when relations are modified.
- ❑ ICs are based upon the semantics of the application that is being described in the database relations.
- ❑ A **legal** instance of a relation is one that satisfies all specified ICs.
  - DBMS should not allow illegal instances.
  - Avoids data entry errors, too!

# Constraints on a Single Relation

- ❑ **not null**
- ❑ **primary key**
- ❑ **unique**
- ❑ **check ( $P$ )**, where  $P$  is a predicate

# Not Null Constraint

- ❑ Declare *branch\_name* for *branch* is **not null**  
*branch\_name* **char(15) not null**

- ❑ Declare the **domain** *Dollars* to be **not null**

**create domain** *Dollars* **numeric(12,2) not null**

# The Unique Constraint

- ❑ **unique** (  $A_1, A_2, \dots, A_m$  )
- ❑ The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key.
- ❑ Candidate keys are permitted to be **null** (in contrast to primary keys).

# Primary Key Constraints



- ❑ A set of attribute is a **key** for a relation if :
  1. No two tuples can have same values in all these attributes, and
  2. This is not true for any subset of the key.
- ❑ Part 2 false?     A **superkey**
- ❑ If there's more than one key for a relation (i.e., **candidate keys**), one is chosen as the **primary key**.

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).

- ❑ *sid* is a key. (What about *name*? What about *login*?)
- ❑ {*sid*, *gpa*} is a superkey.

# Primary & Candidate Keys in SQL

- ❑ Primary key specified as the **PRIMARY KEY**
  - Primary key cannot take a null value
- ❑ Candidate keys specified using **UNIQUE**

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     gpa REAL,
     PRIMARY KEY (sid),
     UNIQUE (login) )
```

- ❑ What's the result of executing the following statements?
  - INSERT INTO Students VALUES ('00001', 'Bob', 'bob@comp', 18, 3.2)
  - INSERT INTO Students VALUES ('00001', 'Tom', 'tom@comp', 18, 3.2)
  - INSERT INTO Students VALUES ('00002', 'Bob', 'bob@comp', 18, 3.2)



# Checking arbitrary constraints

- ❑ We want to make sure that every tuple in STUDENTS has a positive gpa at all times.
- ❑ CREATE TABLE STUDENTS  
( ...,  
PRIMARY KEY(...),  
CHECK (*gpa* > 0))
- ❑ As a result, the database will reject an insertion or an update, if the resulting tuple has a 0 or negative gpa.
- ❑ We can write more complex conditions in CHECK (see next).

## Check – Example 2

- ❑ We already have a table CLUB, recording the GPA of all members in the dancing club.
- ❑ We want to create a table  
**PANEL**(*stu-id*, *major*)  
where each tuple corresponds to a member in the panel of the club.
- ❑ We require that every panel member should have a GPA **at least 1.7**.
- ❑ We can ensure this with CHECK.

CLUB

<i>stu-id</i>	<i>gpa</i>
1	3
2	1.8
3	1.7
4	1.2
5	1.2

PANEL

<i>stu-id</i>	<i>major</i>
1	EE
3	CS

## Check – Example 2

- ❑ We require that every panel member should have a GPA **at least 1.7**.

- ❑ CREATE TABLE PANEL  
(*stu-id* INTEGER,  
*major* CHAR(20),  
PRIMARY KEY (*stu-id*),  
**CHECK** (*stu-id* IN  
(SELECT *stu-id* FROM CLUB  
WHERE *gpa* >= 1.7)))

CLUB

<i>stu-id</i>	<i>gpa</i>
1	3
2	1.8
3	1.7
4	1.2
5	1.2

PANEL

<i>stu-id</i>	<i>major</i>
1	EE
3	CS

- ❑ The database will check the condition whenever
  - there is an insertion/update on PANEL
- ❑ *Note: Oracle does not allow CHECK conditions to use subqueries.*

# The check clause

- ❑ **check** (P)

where P is a predicate

- ❑ The **check** clause in SQL-92 permits **domains** to be restricted:

- Use **check** clause to ensure that an `hourly_wage` domain allows only values greater than a specified value.

```
create domain hourly_wage numeric(5,2)  
           constraint value_test check(value >= 4.00)
```

- The domain has a constraint that ensures that the `hourly_wage` is greater than 4.00
- The clause **constraint** *value\_test* is optional; useful to indicate which constraint an update violated.


# Referential Integrity

- ❑ Example: If “Perryridge” is a branch name appearing in one of the tuples in the *account* relation, then there exists a tuple in the *branch* relation for branch “Perryridge”.
- ❑ Primary and candidate keys and foreign keys can be specified as part of the SQL **create table** statement:
  - The **primary key** clause lists attributes that comprise the primary key.
  - The **unique key** clause lists attributes that comprise a candidate key.
  - The **foreign key** clause lists the attributes that comprise the foreign key and the name of the relation referenced by the foreign key.
    - ▶ By default, a foreign key references the primary key attributes of the referenced table.

# Foreign Key

- ❑ **Foreign key**: Set of attributes in one relation that is used to `refer' to a tuple in another relation (can be itself).
  - Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - E.g., *sid* is a foreign key referring to **Students**

Students(sid: string, name: string, login: string, age: integer, gpa: real)  
Enrolled(sid: string, cid: string, grade: string)



# Foreign Key – Why do we need it?

- ❑ Consider Students and Enrolled:

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❑ Assume that we want to insert a tuple ('50000', 'CS160', 'A') into Enrolled.
- ❑ Before we do so, we may want to make sure there exists a student in Students with *sid* = '50000'.
- ❑ Foreign key is used to achieve this
  - If every *sid* in Enrolled exists in Students, referential integrity is achieved.

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

# Foreign Keys in SQL

- ❑ Only students listed in the *Students* relation should be allowed to **enroll** for courses.
- ❑ But some tuples in *Students* may not be referenced.


```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid, cid),  
FOREIGN KEY (sid) REFERENCES Students )
```

## Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

## Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8






# Foreign Keys in SQL

- ❑ Attribute names can be different.
  - Specify the attribute name explicitly in the definition.

Students(sid: string, name: string, login: string, age: integer, gpa: real)  
Enrolled2(stuid: string, cid: string, grade: string)



```
CREATE TABLE Enrolled2
  (stuid CHAR(20), cid CHAR(20), grade CHAR(2),
  PRIMARY KEY (stuid, cid),
  FOREIGN KEY (stuid) REFERENCES Students (sid) )
```

# Foreign Keys in SQL

- Foreign keys can refer to the same relation. E.g.,



Students2(sid:string, name:string, login:string, age:integer, gpa:real, partner:string)

sid	name	login	age	gpa	partner
53666	John	john@cs	18	2.7	53668
53668	Smith	smith@cs	18	3.8	53666
53650	Smith	smith@ee	19	3.3	NULL

# Foreign Keys in SQL

- ❑ Foreign keys can refer to the same relation. E.g.,



Students2(sid:string, name:string, login:string, age:integer, gpa:real, partner:string)

```
CREATE TABLE Students2
  (sid CHAR(20), name CHAR(20), login CHAR(10),
   age INTEGER, gpa REAL, partner CHAR(20),
   PRIMARY KEY (sid),
   FOREIGN KEY (partner) REFERENCES Students2 (sid) )
```

- If a student has no partner, this field can be **NULL** (a special keyword in SQL denoting *'unknown'* or *'inapplicable'*).
- NULL is allowed in non-primary keys, including foreign keys.

# Enforcing Referential Integrity

- ❑ Consider *Students* and *Enrolled* in the example; *sid* in *Enrolled* is a foreign key that references *Students*.
- ❑ What should be done if an *Enrolled* tuple with a non-existent student id is inserted?  
**Reject it!**
- ❑ What should be done if a *Students* tuple is deleted?
  1. Disallow deletion of a *Students* tuple that is referred to.
  2. Also delete all *Enrolled* tuples that refer to it.
  3. Set *sid* in *Enrolled* tuples that refer to it to a *default sid*.
  4. Another possible option: Set *sid* in *Enrolled* tuples that refer to it to a special value *NULL*.
    - In this example, cannot do it here because *sid* is part of primary key of *Enrolled*.  
→ ***Null is allowed in a foreign key field but not in a primary key field!***
- ❑ Similar if primary key of *Students* tuple is updated.

# Referential Integrity in SQL

❑ SQL/92 and SQL:1999 support all 4 options on deletes and updates.

- Default is **NO ACTION**  
(*delete/update is rejected*)
- **CASCADE** (also delete all tuples that refer to deleted tuple)
- **SET NULL / SET DEFAULT**  
(sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

# Referential Integrity – Example

## Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

## Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❑ What happens if the table Students is updated as follows:

- Delete the tuple with sid = 53666
- Insert a tuple with sid = 53600?
- Update the tuple with sid=53650 → 53700?

# Assertions


- ❑ An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- ❑ Domain constraints and referential-integrity constraints are special forms of assertions.
- ❑ An assertion in SQL takes the form  
**create assertion** <assertion-name> **check** <predicate>
- ❑ When an assertion is made, the system tests it for validity, and tests it again on every update that may violate the assertion
  - This testing may introduce a **significant** amount of **overhead**; hence assertions should be used with great care.
- ❑ SQL does not provide  
for all  $X, P(X)$   
above is achieved by using  
not exists  $X$  such that not  $P(X)$

# Assertion Example

- For each tuple in the *student* relation, the value of the attribute *tot\_cred* must equal the sum of credits of courses that the student has completed successfully.

```
create assertion credits_earned_constraint check  
  (not exists (select ID  
    from student  
    where tot_cred <> (select sum(credits)  
      from takes natural join course  
      where student.ID= takes.ID  
      and grade is not null and grade <> 'F' ))
```





# Triggers

---

# Overview

- ❑ We have learned to write the following constraints.
  - Primary/candidate key (unique)
  - Foreign key
  - CHECK
  - assertion
- ❑ Next, we will discuss another powerful mechanism for writing constraints.
- ❑ Trigger.
- ❑ A trigger can be regarded as a procedure **automatically executed** by the database, **whenever a certain table is modified**.

# Triggers

- ❑ Two requirements of designing a trigger:
  - Specify when a trigger is to be executed.
    - ▶ An event: causes the trigger to be checked
    - ▶ A condition: must be satisfied for trigger execution to proceed
  - Specify the actions to be taken when the trigger executes.

# Triggers

- ❑ Referential integrity on the *time\_slot\_id* attribute of the *section* relation

**create trigger** timeslot\_check1 **after insert** on section

**referencing new row as** nrow

**for each row**



iterate over each inserted row

**when** (nrow.time\_slot\_id **not in** (

**select** time\_slot\_id

**from** time\_slot))




*time\_slot\_id* not  
present in *time\_slot*.

**begin**

**rollback**

**end;**



any transaction that violates  
the referential integrity  
constraint gets rolled back

# Triggers

- ❑ Referential integrity on the *time\_slot\_id* attribute of the *section* relation

**create trigger** timeslot\_check2 **after delete** on *time\_slot*

**referencing old row as** orow

**for each row**

**when** (orow.time\_slot\_id **not in** (

**select** time\_slot\_id

**from** time\_slot)

**and** orow.time\_slot\_id **in** (

**select** time\_slot\_id

**from** section))


**begin**

**rollback**

**end;**



last tuple for *time\_slot\_id*  
deleted from *time\_slot*



and *time\_slot\_id* still  
referenced from *section*

# Triggers

- ❑ Handle grade corrections that change a successful completion grade from a fail grade

**create trigger** credits\_earned **after update of** takes **on** (grade)

**referencing new row as** nrow

**referencing old row as** orow

**for each row**

**when** nrow.grade  $\neq$  'F' **and** nrow.grade **is not null**

**and** (orow.grade = 'F' **or** orow.grade **is null**)

**begin atomic**

**update** student

**set** tot\_cred= tot\_cred+


(**select** credits

**from** course


**where** course.course id= nrow.course id)

**where** student.id = nrow.id;

**end;**



Complete a previously failed course



# Authorization

---

# Authorization Specification in SQL

- ❑ Authorizations on data include:
  - Authorization to read data.
  - Authorization to insert new data.
  - Authorization to update data.
  - Authorization to delete data.
- ❑ Each of these types of authorizations is called a **privilege**.



# Authorization Specification in SQL

- ❑ The **grant** statement is used to confer authorization
  - grant** <privilege list>
  - on** <relation name or view name>
  - to** <user list>
- ❑ <user list> is:
  - a user-id
  - **public**, which allows all valid users the privilege granted
  - A role (more on this later)
- ❑ Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- ❑ The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

# Privileges in SQL

- ❑ **select:** allows **read** access to relation, or the ability to query using the view

- Example: grant users  $U_1$ ,  $U_2$ , and  $U_3$  **select** authorization on the *branch* relation:

**grant select on *branch* to  $U_1$ ,  $U_2$ ,  $U_3$**

- ❑ **insert:** the ability to insert tuples.
- ❑ **update:** the ability to update using the SQL update statement.
- ❑ **delete:** the ability to delete tuples.
- ❑ **all privileges:** used as a short form for all the allowable privileges.

# Revoking Authorization in SQL


**revoke** <privilege list>

**on** <relation name or view name> **from** <user list>

- ❑ **revoke select on** *branch* **from**  $U_1, U_2, U_3$
- ❑ <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- ❑ If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- ❑ If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the **revoke** statement
- ❑ All privileges that depend on the privilege being revoked are also revoked.

# Authorization on Views

- ❑ **create view** *geo\_instructor* **as**  
    **(select** \*  
    **from** *instructor*  
    **where** *dept\_name* = 'Geology');
- ❑ **grant select on** *geo\_instructor* **to** *staff*
- ❑ Suppose that a staff member issues
  - **select** \*  
    **from** *geo\_instructor*;




# Indexes

---

# Indexes

- Primary mechanism to get improved performance on a database
- Persistent data structure, stored in database
- Many interesting implementation issues
  - But we are focusing on user/application perspective


# Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...
	...	...	...

# Functionality




T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...
	...	...	...



# Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...
	...	...	...

# Utility

- Index = difference between full table scans and immediate location of tuples
  - \* Orders of magnitude performance difference
- Underlying data structures
  - Balanced trees (B trees, B+ trees)
  - Hash tables

```
Select sName  
From Student  
where sID = 18942
```

- ❑ Many DBMS's build indexes **automatically** on **PRIMARY KEY** (and sometimes **UNIQUE**) attributes

# Picking which indices to create

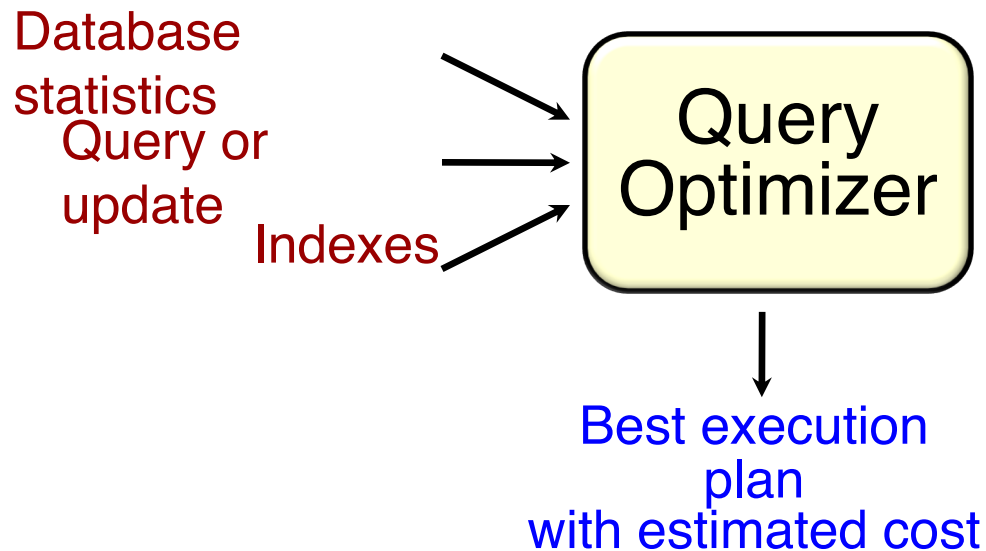
## Benefit of an index depends on:

- Size of table (and possibly layout)
- Data distributions
- Query vs. update load

# “Physical design advisors”

**Input:** database (statistics) and workload

**Output:** recommended indexes



# SQL Syntax

Create Index **IndexName** on T(A)

Create Index **IndexName** on T(A1,A2,...,An)

Create Unique Index **IndexName** on T(A)

Drop Index **IndexName**