


Chapter 4: Advanced SQL

Chapter 4: Advanced SQL

- ❑ Data Types
- ❑ Integrity Constraints
- ❑ Trigger
- ❑ Authorization
- ❑ Index



Data Types

Built-in Data Types in SQL

- ❑ **date:** Dates, containing a (4 digit) year, month and date
 - Example: **date** '2005-7-27'
- ❑ **time:** Time of day, in hours, minutes and seconds.
 - Example: **time** '09:00:30' **time** '09:00:30.75'
- ❑ **timestamp:** date plus time of day
 - Example: **timestamp** '2005-7-27 09:00:30.75'
- ❑ **interval:** period of time
 - Example: **interval** '1' day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values

Build-in Data Types in SQL (Cont.)

- ❑ Can extract values of individual fields from date/time/timestamp
 - **extract** (*field* from *d*)
 - ▶ *field*: can be one of year, month, day, hour, minute, or second.
 - ▶ *d*: date or time value
 - Example: `SELECT EXTRACT(YEAR FROM '2019-07-22')`
- ❑ Can cast string types to date/time/timestamp
 - Example: **cast** <string-valued-expression> **as date**
 - Example: **cast** <string-valued-expression> **as time**

User-Defined Types

- ❑ **create type** construct in SQL creates user-defined type

create type *Dollars* **as numeric (12,2) final**

- **final**: no meaning here, required by the SQL:1999

create table department

(dept name **varchar** (20),

building **varchar** (15),

budget *Dollars*);

- ❑ **create domain** construct in SQL-92 creates similar user-defined domain types

create domain *person_name* **char(20) not null**


- ❑ Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.
- ❑ Domains are not strongly typed. Values of one domain type can be assigned to values of another domain type as long as the underlying types are compatible

Domain Constraints

- ❑ **Domain constraints** are the most elementary form of integrity constraint.
 - They test values inserted in the database, and test queries to ensure that the comparisons make sense.
- ❑ New domains can be created from existing data types
 - Example: **create domain Dollars numeric(12, 2)**
create domain Pounds numeric(12, 2)
- ❑ We cannot assign or compare a value of **type** Dollars to a value of **type** Pounds.
 - However, we can convert type as below
(**cast** *r.A as Pounds*)
(Should also multiply by the dollar-to-pound conversion-rate)

Large-Object Types

- ❑ Large objects (photos, videos, CAD files, etc.) are stored as a *large object*.
 - **blob**: binary large object
 - ▶ object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object
 - ▶ object is a large collection of character data
- book_review* **clob**(10KB)
- image* **blob**(10MB)
- movie* **blob**(2GB)
- When a query returns a large object, a pointer is returned rather than the large object itself.



Integrity Constraints

Integrity Constraints

- ❑ **Integrity Constraints (ICs):** conditions that must be true for **any** instance of the database; e.g., **data type**.
 - An instructor name cannot be null
 - A salary of a bank employee must be at least \$4.00 an hour
 - No two instructors can have the same instructor ID
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- ❑ ICs are based upon the semantics of the application that is being described in the database relations.
- ❑ A **legal** instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
 - Avoids data entry errors, too!

Constraints on a Single Relation

- ❑ **not null**
- ❑ **primary key**
- ❑ **unique**
- ❑ **check (P)**, where P is a predicate

Not Null Constraint

- ❑ Declare *branch_name* for *branch* is **not null**
branch_name **char(15) not null**

- ❑ Declare the **domain** *Dollars* to be **not null**

create domain *Dollars* **numeric(12,2) not null**

The Unique Constraint

- ❑ **unique** (A_1, A_2, \dots, A_m)
- ❑ The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
- ❑ Candidate keys are permitted to be **null** (in contrast to primary keys).

Primary Key Constraints



- ❑ A set of attribute is a key for a relation if :
 1. No two tuples can have same values in all these attributes, and
 2. This is not true for any subset of the key.
- ❑ Part 2 false? A superkey
- ❑ If there's more than one key for a relation (i.e., candidate keys), one is chosen as the primary key.

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).

- ❑ *sid* is a key. (What about *name*? What about *login*?)
- ❑ {*sid*, *gpa*} is a superkey.

Primary & Candidate Keys in SQL

- ❑ Primary key specified as the **PRIMARY KEY**
- ❑ Candidate keys specified using **UNIQUE**

```
CREATE TABLE Students
    (sid CHAR(20),
     name CHAR(20),
     login CHAR(10),
     age INTEGER,
     gpa REAL,
     PRIMARY KEY (sid),
     UNIQUE (login) )
```

- ❑ What's the result of executing the following statements?
 - INSERT INTO Students VALUES ('00001', 'Bob', 'bob@comp', 18, 3.2)
 - INSERT INTO Students VALUES ('00001', 'Tom', 'tom@comp', 18, 3.2)
 - INSERT INTO Students VALUES ('00002', 'Bob', 'bob@comp', 18, 3.2)

Checking arbitrary constraints

- ❑ We want to make sure that every tuple in STUDENTS has a positive gpa at all times.
- ❑ CREATE TABLE STUDENTS
(...,
PRIMARY KEY(...),
CHECK (*gpa* > 0))
- ❑ As a result, the database will reject an insertion or an update, if the resulting tuple has a 0 or negative gpa.
- ❑ We can write more complex conditions in CHECK (see next).

Check – Example 2

- ❑ We already have a table CLUB, recording the GPA of all members in the dancing club.
- ❑ We want to create a table
PANEL(stu-id, major)
where each tuple corresponds to a member in the panel of the club.
- ❑ We require that every panel member should have a GPA *at least 1.7*.
- ❑ We can ensure this with CHECK.

CLUB

<i>stu-id</i>	<i>gpa</i>
1	3
2	1.8
3	1.7
4	1.2
5	1.2

PANEL

<i>stu-id</i>	<i>major</i>
1	EE
3	CS

Check – Example 2

- ❑ We require that every panel member should have a GPA **at least 1.7**.

- ❑ CREATE TABLE PANEL
(*stu-id* INTEGER,
 major CHAR(20),
 PRIMARY KEY (*stu-id*),
 CHECK (*stu-id* IN
 (SELECT *stu-id* FROM CLUB
 WHERE *gpa* >= 1.7)))

CLUB

<i>stu-id</i>	<i>gpa</i>
1	3
2	1.8
3	1.7
4	1.2
5	1.2

PANEL

<i>stu-id</i>	<i>major</i>
1	EE
3	CS

- ❑ The database will check the condition whenever
 - there is an insertion/update on PANEL
- ❑ *Note: Oracle does not allow CHECK conditions to use subqueries.*

The check clause

- ❑ **check** (P)

where P is a predicate

- ❑ The **check** clause in SQL-92 permits **domains** to be restricted:

- Use **check** clause to ensure that an `hourly_wage` domain allows only values greater than a specified value.

```
create domain hourly_wage numeric(5,2)  
               constraint value_test check(value >= 4.00)
```

- The domain has a constraint that ensures that the `hourly_wage` is greater than 4.00
- The clause **constraint** *value_test* is optional; useful to indicate which constraint an update violated.


Referential Integrity

- ❑ Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Perryridge” is a branch name appearing in one of the tuples in the *account* relation, then there exists a tuple in the *branch* relation for branch “Perryridge”.
- ❑ Primary and candidate keys and foreign keys can be specified as part of the SQL **create table** statement:
 - The **primary key** clause lists attributes that comprise the primary key.
 - The **unique key** clause lists attributes that comprise a candidate key.
 - The **foreign key** clause lists the attributes that comprise the foreign key and the name of the relation referenced by the foreign key. By default, a foreign key references the primary key attributes of the referenced table.

Foreign Key

- ❑ **Foreign key**: Set of attributes in one relation that is used to `refer' to a tuple in another relation (can be itself).
 - Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - E.g., *sid* is a foreign key referring to **Students**

Students(sid: string, name: string, login: string, age: integer, gpa: real)
Enrolled(sid: string, cid: string, grade: string)



Foreign Key – Why do we need it?

- ❑ Consider Students and Enrolled:

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- ❑ Assume that we want to insert a tuple ('50000', 'CS160', 'A') into Enrolled.
- ❑ Before we do so, we may want to make sure there exists a student in Students with *sid* = '50000'.
- ❑ Foreign key is used to achieve this
 - If every sid in Enrolled exists in Students, referential integrity is achieved.

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Foreign Keys in SQL

- ❑ Only students listed in the Students relation should be allowed to **enroll** for courses.
- ❑ But some tuples in *Students* may not be referenced.


```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid, cid),  
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



Foreign Keys in SQL

- ❑ Attribute names can be different.
 - Specify the attribute name explicitly in the definition.

Students(sid: string, name: string, login: string, age: integer, gpa: real)
Enrolled2(stuid: string, cid: string, grade: string)



```
CREATE TABLE Enrolled2
  (stuid CHAR(20), cid CHAR(20), grade CHAR(2),
  PRIMARY KEY (stuid, cid),
  FOREIGN KEY (stuid) REFERENCES Students (sid) )
```


Foreign Keys in SQL

- Foreign keys can refer to the same relation. E.g.,



Students2(sid:string, name:string, login:string, age:integer, gpa:real, partner:string)

sid	name	login	age	gpa	partner
53666	John	john@cs	18	2.7	53668
53668	Smith	smith@cs	18	3.8	53666
53650	Smith	smith@ee	19	3.3	NULL

Foreign Keys in SQL

- ❑ Foreign keys can refer to the same relation. E.g.,



Students2(sid:string, name:string, login:string, age:integer, gpa:real, partner:string)

```
CREATE TABLE Students2
  (sid CHAR(20), name CHAR(20), login CHAR(10),
   age INTEGER, gpa REAL, partner CHAR(20),
   PRIMARY KEY (sid),
   FOREIGN KEY (partner) REFERENCES Students2 (sid) )
```

- If a student has no partner, this field can be **NULL** (a special keyword in SQL denoting *'unknown'* or *'inapplicable'*).
- NULL is allowed in non-primary keys, including foreign keys.

Enforcing Referential Integrity

- ❑ Consider **Students** and **Enrolled** in the example; *sid* in Enrolled is a foreign key that references Students.
- ❑ What should be done if an Enrolled tuple with a non-existent student id is inserted?
Reject it!
- ❑ What should be done if a Students tuple is deleted?
 1. Disallow deletion of a Students tuple that is referred to.
 2. Also delete all Enrolled tuples that refer to it.
 3. Set *sid* in Enrolled tuples that refer to it to a **default sid**.
 4. Another possible option: Set *sid* in Enrolled tuples that refer to it to a special value **NULL**.
 - In this example, cannot do it here because *sid* is part of primary key of Enrolled.

→ ***Null is allowed in a foreign key field but not in a primary key field!***
- ❑ Similar if primary key of Students tuple is updated.

Referential Integrity in SQL

❑ SQL/92 and SQL:1999 support all 4 options on deletes and updates.

- Default is **NO ACTION**
(*delete/update is rejected*)
- **CASCADE** (also delete all tuples that refer to deleted tuple)
- **SET NULL / SET DEFAULT**
(sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '53688',
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

Referential Integrity – Example

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❑ What happens if the table Students is updated as follows:

- Delete the tuple with sid = 53666
- Insert a tuple with sid = 53600?
- Update the tuple with sid=53650 → 53700?


Assertions

- ❑ An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- ❑ Domain constraints and referential-integrity constraints are special forms of assertions.
- ❑ An assertion in SQL takes the form
create assertion <assertion-name> **check** <predicate>
- ❑ When an assertion is made, the system tests it for validity, and tests it again on every update that may violate the assertion
 - This testing may introduce a **significant** amount of **overhead**; hence assertions should be used with great care.
- ❑ SQL does not provide
for all $X, P(X)$
above is achieved by using
not exists X such that not $P(X)$

Assertion Example

- ❑ For each tuple in the student relation, the value of the attribute tot cred must equal the sum of credits of courses that the student has completed successfully.

```
create assertion credits_earned_constraint check  
  (not exists (select ID  
    from student  
    where tot_cred <> (select sum(credits)  
      from takes natural join course  
      where student.ID= takes.ID  
      and grade is not null and grade <> 'F' ))
```



Triggers

Triggers

- ❑ “Event-Condition-Action Rules”
- ❑ When *event* occurs, check *condition*; if true, do *action*
 - 1) Move monitoring logic from apps into DBMS
 - 2) Enforce constraints
 - Beyond what constraint system supports
 - Automatic constraint “repair”
- ❑ Implementations vary significantly

Overview

- ❑ We have learned to write the following constraints.
 - Primary/candidate key (unique)
 - Foreign key
 - CHECK
 - assertion
- ❑ Next, we will discuss another powerful mechanism for writing constraints.
- ❑ **Trigger.**
- ❑ A trigger can be regarded as a procedure **automatically executed** by the database, **whenever a certain table is modified**.

Triggers

- ❑ Two requirements of designing a trigger:
 - Specify when a trigger is to be executed.
 - ▶ An event: causes the trigger to be checked
 - ▶ A condition: must be satisfied for trigger execution to proceed
 - Specify the actions to be taken when the trigger executes.

Triggers

- ❑ Referential integrity on the *time_slot_id* attribute of the *section* relation

create trigger timeslot_check1 **after insert** on section

referencing new row as nrow

for each row



iterate over each inserted row

when (nrow.time_slot_id **not in** (

select time_slot_id

from time_slot))




time_slot_id not
present in *time_slot*.

begin

rollback

end;



any transaction that violates
the referential integrity
constraint gets rolled back

Triggers

- ❑ Referential integrity on the *time_slot_id* attribute of the *section* relation

create trigger timeslot_check2 **after delete** **on** time_slot

referencing old row as orow

for each row

when (orow.time_slot_id **not in** (

select time_slot_id

from time_slot)

and orow.time_slot_id **in** (

select time_slot_id

from section))


begin

rollback

end;



last tuple for *time_slot_id*
deleted from *time_slot*



and *time_slot_id* still
referenced from *section*

Triggers

- ❑ Handle grade corrections that change a successful completion grade from a fail grade

create trigger credits_earned **after update of** takes **on** (grade)

referencing new row as nrow

referencing old row as orow

for each row

when nrow.grade \neq 'F' **and** nrow.grade **is not null**

and (orow.grade = 'F' **or** orow.grade **is null**)

begin atomic

update student

set tot cred = tot cred +


(**select** credits

from course


where course.course id = nrow.course id)

where student.id = nrow.id;

end;



Complete a previously failed course



Authorization

Authorization Specification in SQL

- ❑ Authorizations on data include:
 - Authorization to read data.
 - Authorization to insert new data.
 - Authorization to update data.
 - Authorization to delete data.
- ❑ Each of these types of authorizations is called a **privilege**.

Authorization Specification in SQL

- ❑ The **grant** statement is used to confer authorization
 - grant** <privilege list>
 - on** <relation name or view name>
 - to** <user list>
- ❑ <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this later)
- ❑ Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- ❑ The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

Privileges in SQL

- ❑ **select:** allows **read** access to relation, or the ability to query using the view

- Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:

grant select on *instructor* to U_1 , U_2 , U_3

- ❑ **insert:** the ability to insert tuples.
- ❑ **update:** the ability to update using the SQL update statement.
- ❑ **delete:** the ability to delete tuples.
- ❑ **all privileges:** used as a short form for all the allowable privileges.

Revoking Authorization in SQL


revoke <privilege list>

on <relation name or view name> **from** <user list>

- ❑ **revoke select on** *branch* **from** U_1, U_2, U_3
- ❑ <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- ❑ If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- ❑ If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the **revoke** statement
- ❑ All privileges that depend on the privilege being revoked are also revoked.

Authorization on Views

- ❑ **create view** *geo_instructor* **as**
 (**select** *
 from *instructor*
 where *dept_name* = 'Geology');
- ❑ **grant select on** *geo_instructor* **to** *staff*
- ❑ Suppose that a staff member issues
 - **select** *
 from *geo_instructor*;




Indexes

Indexes

- Primary mechanism to get improved performance on a database
- Persistent data structure, stored in database
- Many interesting implementation issues
 - But we are focusing on user/application perspective


Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...


Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...

Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...

Utility

- Index = difference between full table scans and immediate location of tuples
 - * Orders of magnitude performance difference
- Underlying data structures
 - Balanced trees (B trees, B+ trees)
 - Hash tables

```
Select sName  
From Student  
where sID = 18942
```

- ❑ Many DBMS's build indexes **automatically** on **PRIMARY KEY** (and sometimes **UNIQUE**) attributes

```
Select sID  
From Student  
where sName = 'Mary' And GPA >  
3.9
```

```
Select sName, cName  
From Student, Apply  
where Student.sID = Apply.sID
```

Picking which indices to create

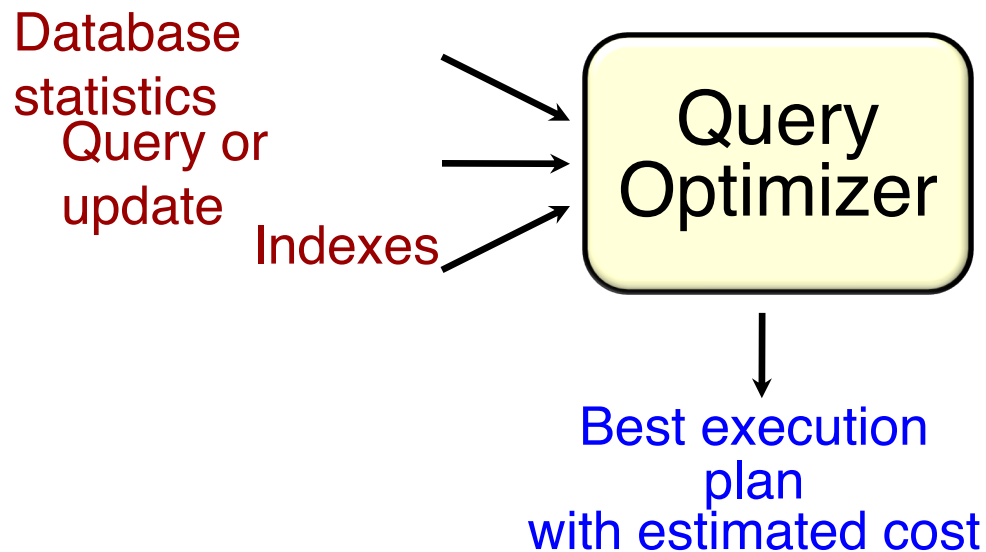
Benefit of an index depends on:

- Size of table (and possibly layout)
- Data distributions
- Query vs. update load

“Physical design advisors”

Input: database (statistics) and workload

Output: recommended indexes



SQL Syntax

Create Index **IndexName** on T(A)

Create Index **IndexName** on T(A1,A2,...,An)

Create Unique Index **IndexName** on T(A)

Drop Index **IndexName**